

## **OBJECTIVES**

- **Define COMMIT and ROLLBACK**
- **To understand the need for concurrency control and backup/recovery**
- **To learn about typical problems that can occur when multiple users process a database concurrently**
- **To understand the use of locking and the problem of deadlock**
- **To learn the difference between optimistic and pessimistic locking**

## **More OBJECTIVES**

- **To know the meaning of an ACID transaction**
- **To learn the four 1992 ANSI standard isolation levels**
- **To know the difference between recovery via reprocessing and recovery via rollback/rollforward**
- **To understand the nature of the tasks required for recovery using rollback/rollforward**

# ***DBMS Transaction Logging***

---

**Start with BACKUP:**

**The DBAs' “bottom line” – never, ever lose data**

So we run backups. How often?

- It depends on the size of the database, the number of users, the frequency of updates, and how critical the database is to the organization.

Some systems are SO CRITICAL that we keep redundant copies up-to-date at all times. Lose one, use the other.

# ***DBMS Transaction Logging***

---

**Redundant copies of a database:**

**Relies on technology called REPLICATION –**

**a topic for a later date...**

# ***DBMS Transaction Logging***

---

**Backups and Replication both rely on**

## **Transaction Logs**

**Every activity against the database that changes the database is written to a log file.**

**Log files reside in memory and are written to disk periodically. (In MySQL, this is called a “flush”.)**

**Transactions are written to the log file first, then written to the database.**

# ***DBMS Transaction Logging***

---

**Types of backups – usually we combine all of these**

**1. Stop all database traffic and back it up quickly**

- “Cold” backup (“full” or “incremental”)
- MySQLDUMP versus a file-level backup (OS level)
- Faster
- Requires downtime

**2. Take the backup while database traffic is in-progress**

- “Warm” backup (“full” or “incremental”)
- Takes longer
- Less inconvenience to user community

## **Types of backups**

### **3. Keep a duplicate up-to-date at all times**

- “Hot standby” via replication from logs

### **4. Full versus Incremental**

- The backup software copies the entire database (“full”) versus copying only rows changed since the last “full” (“incremental”)

# ***DBMS Transaction Logging***

---

**Logs can be used to UNDO a transaction or REDO a transaction.**

**Logs capture an image of the row BEFORE and AFTER it was changed.**

**Logs capture timestamps for every activity.**



# ***DBMS Transaction Logging***


---

**Recovering a database after a crash or corruption:**

- 1. Identify the most recent full backup**
- 2. Identify the latest transaction log file**
- 3. Identify the point-in-time of the failure**
- 4. Restore the most recent full backup**
- 5. Apply transaction logs up to the last commit before the failure occurred**


# *DBMS Transaction Logging*

**Before-image:** a copy of every database record (or page) before it was changed



# *DBMS Transaction Logging*

**After-image:** a copy of every database record (or page) after it was changed



# *DBMS Transaction Logging*

---

- **Logs in memory must be written to disk**
- **A log flush in MySQL**
- **(“redo logs” → “archive logs” in Oracle.)**
- **Log files can be automatically written to disk**
  - On a timer, every nn minutes
  - On a size limit, whenever the log file grows to nn MB or GB
- **Closes out the old log file and opens up a new one**
  - Increments the log number in the filename

# DBMS Transaction Logging

Relative Record Number	Transaction ID	Reverse Pointer	Forward Pointer	Time	Type of Operation	Object	Before-Image	After-Image
1	OT1	0	2	11:42	START			
2	OT1	1	4	11:43	MODIFY	CUST 100	(old value)	(new value)
3	OT2	0	8	11:46	START			
4	OT1	2	5	11:47	MODIFY	SP AA	(old value)	(new value)
5	OT1	4	7	11:47	INSERT	ORDER 11		(value)
6	CT1	0	9	11:48	START			
7	OT1	5	0	11:49	COMMIT			
8	OT2	3	0	11:50	COMMIT			
9	CT1	6	10	11:51	MODIFY	SP BB	(old value)	(new value)
10	CT1	9	0	11:51	COMMIT			

# *Concurrency Control*

---

- **Concurrency control** ensures that one user's work does not inappropriately influence another user's work.
  - No single concurrency control technique is ideal for all circumstances.
  - Trade-offs need to be made between level of protection and throughput.

- A **transaction**, or **logical unit of work (LUW)**, is a series of actions taken against the database that occurs as an atomic unit:
  - Either all actions in a transaction occur or none of them do.

# Transaction Control

## What can happen without transaction control

Before

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

### SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$3200	\$320

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

Action

- START
1. Add new-order data to CUSTOMER.
  2. Add new-order data to SALESPERSON.
  3. Insert new ORDER.
- STOP

After

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400
123	8000	250 Basketballs	\$6500

### SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$9700	\$970

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*



# Transaction Control

## With atomic transaction control

Before

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

### SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$3200	\$320

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

Transaction

Begin Transaction  
Change CUSTOMER data  
Change SALESPERSON data  
Insert ORDER data  
If no errors then  
Commit Transactions  
Else  
Rollback Transaction  
End If

After

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

### SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$3200	\$320

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

# *Transaction Control*


---

- **Concurrent transactions** refer to two or more transactions that appear to users as they are being processed against a database at the same time.
- In reality, CPU can execute only one instruction at a time.
  - **Transactions are interleaved** meaning that the operating system quickly switches CPU services among tasks so that some portion of each of them is carried out in a given interval.
- **Concurrency problems are lost updates and inconsistent reads.**


# Transaction Control

- Sequence of events without conflict

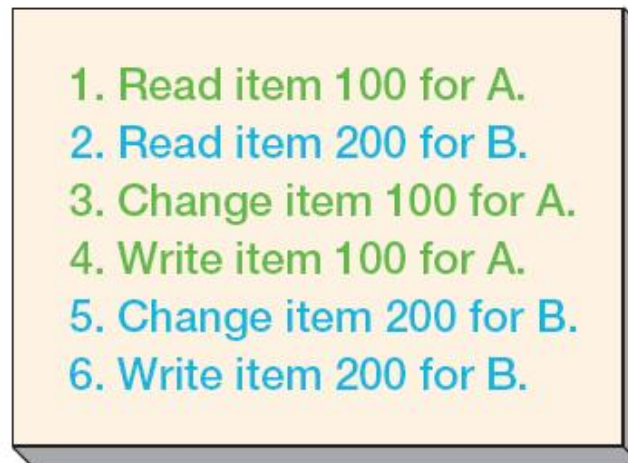
User A



User B



Order of processing at database server



# Transaction Control

- Events with conflict – Lost Update

User A

1. Read item 100  
(item count is 10).
2. Reduce count of items by 5.
3. Write item 100.

User B

1. Read item 100  
(item count is 10).
2. Reduce count of items by 3.
3. Write item 100.

Order of processing at database server

1. Read item 100 (for A).
2. Read item 100 (for B).
3. Set item count to 5 (for A).
4. Write item 100 for A.
5. Set item count to 7 (for B).
6. Write item 100 for B.

Note: The change and write in steps 3 and 4 are lost.

# *Handling Concurrency*

---

- **Resource locking** prevents multiple applications from obtaining copies of the same record when the record is about to be changed.

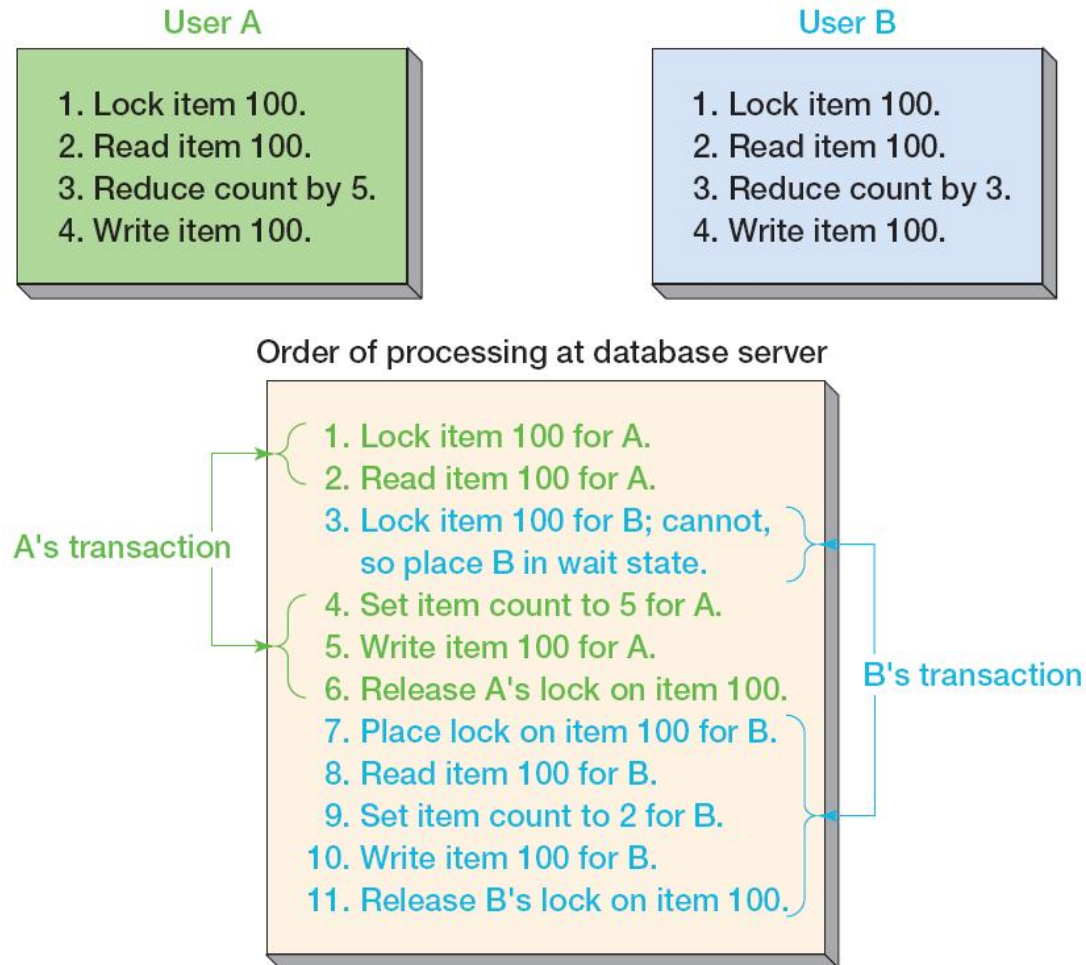
# *Handling Concurrency*

---

- **Implicit locks** are locks placed by the DBMS.
- **Explicit locks** are issued by the application program.
- **Lock granularity** refers to size of a locked resource:
  - Rows, page, table, and database level.
- **Large granularity is easier to manage but frequently causes conflicts.**
- **Types of lock:**
  - An **exclusive lock** prohibits other users from reading the locked resource.
  - A **shared lock** allows other users to read the locked resource, but they cannot update it.

# Handling Concurrency

- Serialized Transactions



# *Handling Concurrency*

---

- **Serializable transactions** refer to two transactions that run concurrently and generate results that are consistent with the results that would have occurred if they had run separately.
- **Two-phased locking** is one of the techniques used to achieve serializability.



- **Two-phased locking:**
  - Transactions obtain locks as necessary (**growing phase**) until all necessary objects are locked.
  - Once any lock is released (**shrinking phase**), subsequent locks will be released until all are released, but the transaction cannot grab more new locks.
- **A special simplified case of two-phased locking:**
  - Locks are obtained throughout the transaction.
  - No lock is released until the COMMIT or ROLLBACK command is issued.
  - This strategy is more restrictive but easier to implement than two-phase locking.

# *Handling Concurrency*

---

- **Deadlock**, or the deadly embrace, occurs when two transactions are each waiting on a resource that the other transaction holds.
- **Preventing deadlock:**
  - Allow users to issue all lock requests up front.
- **Breaking deadlock:**
  - Almost every DBMS has algorithms for detecting deadlock.
  - When deadlock occurs, DBMS aborts one of the transactions (selected randomly ), allowing one transaction to complete, while the aborted transaction rolls back any partially completed work.

# Handling Concurrency

User A

1. Lock paper.
2. Take paper.
3. Lock pencils.

User B

1. Lock pencils.
2. Take pencils.
3. Lock paper.

Order of processing at database server

1. Lock paper for user A.
2. Lock pencils for user B.
3. Process A's requests; write paper record.
4. Process B's requests; write pencil record.
5. Put A in wait state for pencils.
6. Put B in wait state for paper.

\*\* Locked \*\*

# *Handling Concurrency*

---

- **Most application programs do not explicitly declare locks due to the complexities of managing them. An exception: some non-relational DBMS systems.**
- **Instead, they mark *transaction boundaries* and declare the locking behavior they want the DBMS to use.**
  - Transaction boundary markers: START, COMMIT, and ROLLBACK TRANSACTION.

# *Handling Concurrency*

- **Psuedo-Code**

```
BEGIN TRANSACTION

SELECT      PRODUCT.Name, PRODUCT.Quantity
FROM        PRODUCT
WHERE       PRODUCT.Name = 'Pencil'

Set NewQuantity = PRODUCT.Quantity - 5

{process transaction - take exception action if NewQuantity < 0, etc.}

UPDATE      PRODUCT
SET         PRODUCT.Quantity = NewQuantity
WHERE       PRODUCT.Name = 'Pencil'

{continue processing transaction} . . .

IF {transaction has completed normally} THEN

    COMMIT TRANSACTION

ELSE

    ROLLBACK TRANSACTION

END IF

Continue processing other actions not part of this transaction . . .
```

- The acronym **ACID**
  - Transaction management that is **A**tomic, **C**onsistent, **I**solated, and **D**urable.
- **Atomic** means either **ALL** or **NONE** of the database actions within a transaction are committed.
- **Durable** means database committed changes are permanent.

# *Handling Concurrency*

---

- **Consistency** means the DBMS software will not allow any application programs to violate any database constraints. Once committed, the transaction leaves the data in a consistent state.

- **Isolation** means the DBMS software
  - Allows application programmers are able to declare the desired isolation level,
  - So that they can trust that the the DBMS manages locks so as to achieve that level of isolation.



# *Handling Concurrency*

---

- The ANSI SQL-92 standard defines four **transaction isolation** levels:
  - Read uncommitted – allows reads of uncommitted changes ‘i.e. “dirty” reads, unrepeatable reads and phantom reads
  - Read committed – allows unrepeatable reads and phantom reads, but no dirty reads
  - Repeatable read – allows phantom reads, but no unrepeatable reads
  - Serializable – guarantees no read anomalies

# *Handling Concurrency*

		Isolation Level			
		Read Uncommitted	Read Committed	Repeatable Read	Serializable
Problem Type	Dirty Read	Possible	Not Possible	Not Possible	Not Possible
	Nonrepeatable Read	Possible	Possible	Not Possible	Not Possible
	Phantom Read	Possible	Possible	Possible	Not Possible

# *Handling Concurrency*

---

- **Isolation**
- **What's the difference?**
  - The speed at which things get done
  - The frequency and duration of various locks
- **Trade-offs**
  - Faster throughput versus data consistency and reliability