

OBJECTIVES


- **To understand how indexes work**
- **Types of MySQL Indexes**

The concept of an index

- Think of the index at the back of a book
- Speeds up a search
 - Look up a key word in the index
 - Index points to a page number in the book
- Without the index, you would have to scan every page in the book looking for your key word

Consider database server architecture

MySQL
Database Server




Relative Speeds

nanoseconds
microseconds
milliseconds

MySQL Indexes

Data is moved from
disk to memory
one page at a time.

MySQL
Database Server



MySQL Default page size = 16K

To read an entire table, the table is moved into memory one page at a time.

Any select, insert, update, delete requires the affected page where that row resides to be brought into memory (“buffer cache.”)

DML statements modify affected rows on the page in memory. Modified pages are written to disk upon a “commit”.

Moving pages into and out of memory is relatively very time consuming.

MySQL Indexes

MySQL Memory

Server Shared
Storage Engine Shared
Connection Session

Query Cache
Thread Cache

Log Buffer
Buffer Pool

Sort Buffer
Read Buffer
Temp Table

MySQL memory allocation is divided into two categories:

- Global (per instance)

Allocated when the instance starts and freed when the instance is shut down. If it fills up, the OS will swap, DB performance will suffer, and eventually the server will crash.

- Session (per connection)

Allocated per session. Released at session end. Used for query results. Read buffer size is per session. 10MB read buffer * 100 concurrent sessions = 1GB of memory.

Query Cache

- Holds the compiled executable query.
- Queries may be run over and over by an application, changing only host variables.
 - Think of a drop-down list created at web page load that reads many rows.
 - Caching the query avoids the time spent validating, parsing, and compiling the query.
- Deprecated in MySQL 5.7
 - Reduces performance predicatability
 - Doesn't scale with huge workloads on multi-core machines

Thread Cache

- Splits up the execution paths of MySQL into multiple threads, one per user connection, so they run in parallel.

Buffer Pools

- Caching disk I/O
- Can hold objects in memory (smaller tables, indexes)

Session Memory

- Smaller temp tables
- Client-specific memory
 - Table read buffers
 - Sort operations

So why indexes?

We use indexes to minimize page movement (i.e. Minimize “disk I/O”).

Indexes are usually much smaller than the base tables they serve -- less I/O, less buffer space consumed.

Indexes are often pinned in memory.

Simple index structure.

Relative Record Number	Student- Number	Class- Number	Semester
1	200	70	2000S
2	100	30	2001F
3	300	20	2001F
4	200	30	2000S
5	300	70	2000S
6	100	20	2000S

(a) ENROLLMENT Data

Student- Number	Relative Record Number
100	2
100	6
200	1
200	4
300	3
300	5

(b) Index on
StudentNumber

Class- Number	Relative Record Number
20	3
20	6
30	2
30	4
70	1
70	5

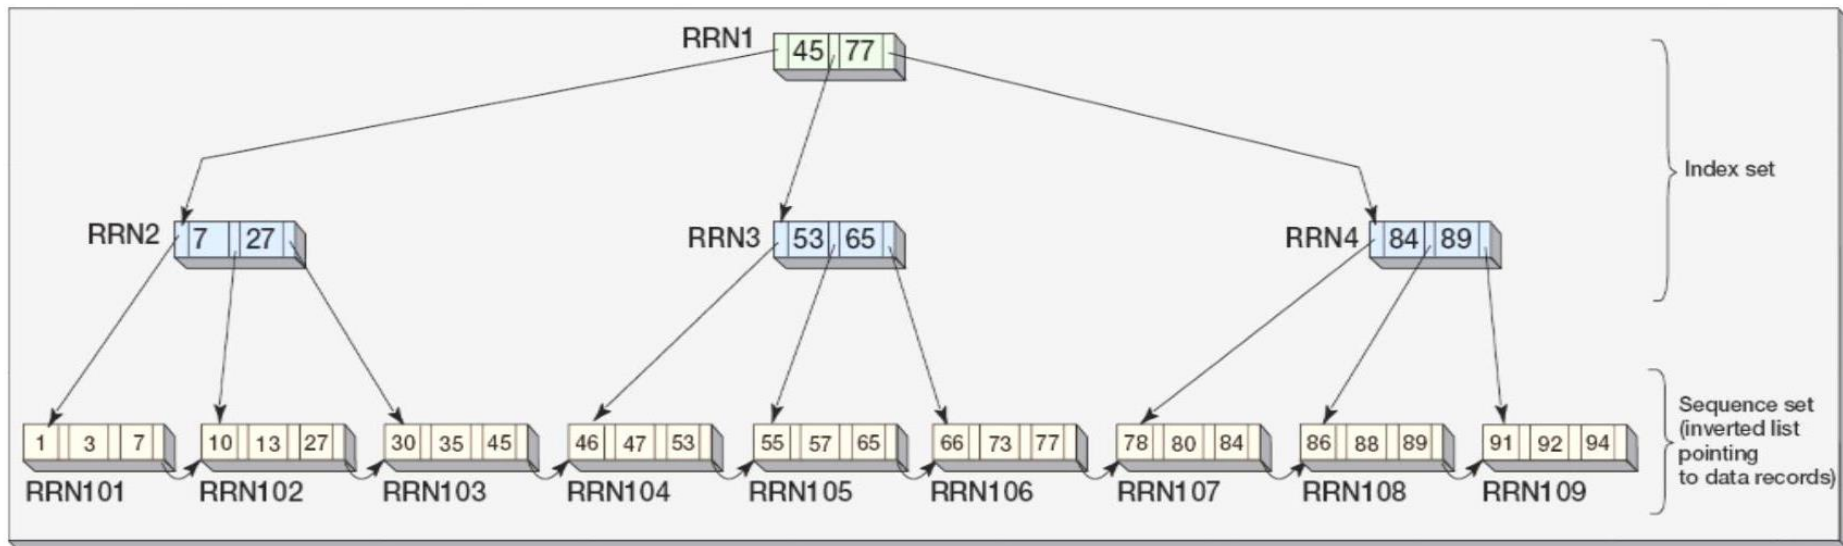
(c) Index on
ClassNumber

MySQL Indexes

MySQL indexes are typically B-Tree, the default.

- B-Tree supports both sequential access of rows AND keyed access.
- “Index Set and Sequence Set.”
- “root” and “leaf” nodes.

General Structure of a Simple B-Tree



- The “B” means “balanced”
- B-Tree indexes balance each leaf page node between half-full and full.
- B-Tree indexes balance the levels of leaf pages.
- As pages get full, B-Tree indexes will “split” to create more room for indexes to grow. The index has pre-allocated overflow space.

- As the overflow space gets full, the “split” leaf nodes get spread out.
- Therefore B-Tree indexes occasionally need to be reorganized.
- Drop the index. Recreate it.

- The default DB engine for MySQL is “innodb”.
- A database or table can be created to use the “Memory” DB engine. All objects are memory resident, asynchronously written to disk.
- The Memory engine supports “hash” indexes. The key is run through a hashing algorithm which calculates a shorter key length, which is then associated with a pointer.

- Primary keys are indexed by default.
- Data rows are written to disk in primary key sequence. This is a “clustered” index.
- Allows for binary search.

- Foreign keys may be indexed.
- Other non-key columns may be indexed. (“Secondary” index.)
- The data rows will not be in secondary key sequence.
- MySQL secondary indexes include the primary key values so that access by secondary keys can leverage the clustered index.

- How do we know what secondary indexes are needed?
- Database designers must carefully assess and analyze the full collection of the organization's queries.
- Secondary indexes are based on query usage patterns.

- Indexes present a trade-off:
 - READs go faster when indexed
 - INSERTs require not only data table updates, but also index updates.
 - For applications that are insert-intensive, indexes can cause significant delays.
 - For small tables, an index can actually hurt performance

- SO –
 - Be very careful about creating indexes
 - Only create a new index when query performance demands it

- **“Ad Hoc”** query access.
 - End-User Reporting tools generate SQL
 - Tableau
 - Business Objects
 - Introduces huge challenges
 - Queries must be monitored for performance
 - Performance bottlenecks must be managed
 - Requires 3rd party tools.
 - WebYog, DataDog for monitoring
 - Queuing - Not in MySQL, but other DBMSs have it

- MySQL allows multi-column indexes.
- MySQL will only use the N^{th} column in the index if the $(N-1)^{\text{th}}$ column is used first.
- If index columns are carefully chosen, a single index can boost performance for many different queries.

- For example:

```
CREATE TABLE test (  
    id INT NOT NULL,  
    last_name CHAR(30) NOT NULL,  
    first_name CHAR(30) NOT NULL,  
    PRIMARY KEY (id),  
    INDEX IX_Name (last_name, first_name)  
);
```


These WILL use `IX_Name (last_name, first_name)`

- `SELECT * FROM test WHERE last_name='Widenius';`
- `SELECT * FROM test WHERE last_name='Widenius' AND first_name='Michael';`
- `SELECT * FROM test WHERE last_name='Widenius' AND (first_name='Michael' OR first_name='Monty');`
- `SELECT * FROM test WHERE last_name='Widenius' AND first_name >='M' AND first_name < 'N';`

- **These WILL NOT use** `IX_Name (last_name, first_name)`
- `SELECT * FROM test WHERE first_name='Michael';`
- `SELECT * FROM test WHERE last_name='Widenius' OR first_name='Michael';`

- Indexes can be included in the TABLE CREATE
- Indexes can be dropped, renamed, added later via a ALTER TABLE

```
ALTER TABLE contacts  
    RENAME INDEX contacts_idx  
    TO    contacts_new_index;
```

```
DROP INDEX contacts_idx  
    ON contacts;
```

MySQL keeps statistics on indexes in the “statistics” table in the “information_schema” table. And, innodb keeps table and index stats in the mysql schema.

The query optimizer uses these statistics as it parses a query and builds the execution plan to decide whether or not to use an index.

If the statistics are “stale”, the optimizer might decide to NOT use an important index.

Therefore, statistics need to be refreshed.

MySQL Indexes

Table 14.3 Columns of innodb_table_stats

Column name	Description
database_name	Database name
table_name	Table name, partition name, or subpartition name
last_update	A timestamp indicating the last time that InnoDB updated this row
n_rows	The number of rows in the table
clustered_index_size	The size of the primary index, in pages
sum_of_other_index_sizes	The total size of other (non-primary) indexes, in pages

Table 14.4 Columns of innodb_index_stats

Column name	Description
database_name	Database name
table_name	Table name, partition name, or subpartition name
index_name	Index name
last_update	A timestamp indicating the last time that InnoDB updated this row
stat_name	The name of the statistic, whose value is reported in the <code>stat_value</code> column
stat_value	The value of the statistic that is named in <code>stat_name</code> column
sample_size	The number of pages sampled for the estimate provided in the <code>stat_value</code> column
stat_description	Description of the statistic that is named in the <code>stat_name</code> column

For example, if massive inserts are done that shifts the cardinality (number of unique values in the domain of an indexed column), the optimizer may choose NOT to use the index.

Changes to > 10% of the rows triggers a stats recalc automatically

Solutions:

1. Optimize table – rebuilds the table and all indexes.
2. Analyze table – forces a stats recalc